

FinalProject_Report

Kriti Gupta

29/09/2020

1.Introduction

The purpose for this project is creating a recommender system using MovieLens dataset.

The version of movielens dataset used for this final assignment contains approximately 10 Millions of movies ratings, divided in 9 Millions for training and one Million for validation. It is a small subset of a much larger (and famous) dataset with several millions of ratings. After a initial data exploration, the recommender systems built on this dataset are evaluated and choosen based on the RMSE - Root Mean Squared Error that should be at least lower than **0.89999**.

#Installing essential packages

First the working environment is set up by installing essential packages:

Install all needed libraries if it is not present

```
if(!require(tidyverse)) install.packages("tidyverse")
if(!require(kableExtra)) install.packages("kableExtra")
if(!require(tidyr)) install.packages("tidyr")
if(!require(tidyverse)) install.packages("tidyverse")
if(!require(stringr)) install.packages("stringr")
if(!require(forcats)) install.packages("forcats")
if(!require(ggplot2)) install.packages("ggplot2")
if(!require(data.table)) install.packages("data.table")
```

Loading all needed libraries

```
library(dplyr)
library(tidyverse)
library(kableExtra)
library(tidyr)
library(stringr)
library(forcats)
library(ggplot2)
library(lubridate)
library(caret)
library(tinytex)
library(data.table)
```

2. Data downloading and preparation


Now we split the MovieLens dataset into Training (*edx*) and Validation (*validation*) sets. The Validation set will be 10% of MovieLens data.

3.Data exploration

##3.1 Overall profile of the dataset Let's first have a general overview of the dataset:

```
head(edx)
```

```
##      userId movieId rating timestamp                title
## 1         1      122      5 838985046          Boomerang (1992)
## 2         1      185      5 838983525            Net, The (1995)
## 4         1      292      5 838983421            Outbreak (1995)
## 5         1      316      5 838983392            Stargate (1994)
## 6         1      329      5 838983392 Star Trek: Generations (1994)
## 7         1      355      5 838984474    Flintstones, The (1994)
##                                     genres
## 1                                Comedy|Romance
## 2                        Action|Crime|Thriller
## 4    Action|Drama|Sci-Fi|Thriller
## 5                Action|Adventure|Sci-Fi
## 6    Action|Adventure|Drama|Sci-Fi
## 7                Children|Comedy|Fantasy
```



We can see the overall distribution of all of the ratings. It is skewed to the right. All half stars are less frequent than full stars. A red dashed line of the overall average rating is also plotted here as a reference.

```
dim(edx) # 9000055      6
n_distinct(edx$movieId) # 10677
n_distinct(edx$title) # 10676: there might be movies of different IDs with the same title
n_distinct(edx$userId) # 69878
n_distinct(edx$movieId)*n_distinct(edx$userId) # 746087406
n_distinct(edx$movieId)*n_distinct(edx$userId)/dim(edx)[1] # 83
```

3.2 Extracting age of movies


```
edx_1 <- edx %>% mutate(year Rated = year(as_datetime(timestamp)))
# extract the release year of the movie
# edx_1 has year Rated, year Released, age_at_rating, and titles without year information
edx_1 <- edx_1 %>% mutate(title = str_replace(title, "^(.+)\s\\((\\d{4})\\)$", "\\1_\\2" )) %>%
  separate(title, c("title", "year Released"), "_") %>%
  select(-timestamp)
edx_1 <- edx_1 %>% mutate(age_at_rating = as.numeric(year Rated) - as.numeric(year Released))
head(edx_1)
```


##	userId	movieId	rating	title	year Released
## 1	1	122	5	Boomerang	1992
## 2	1	185	5	Net, The	1995

##	3	1	292	5	Outbreak	1995
##	4	1	316	5	Stargate	1994
##	5	1	329	5	Star Trek: Generations	1994
##	6	1	355	5	Flintstones, The	1994
##					genres	year Rated age_at_rating
##	1				Comedy Romance	1996 4
##	2				Action Crime Thriller	1996 1
##	3				Action Drama Sci-Fi Thriller	1996 1
##	4				Action Adventure Sci-Fi	1996 2
##	5				Action Adventure Drama Sci-Fi	1996 2
##	6				Children Comedy Fantasy	1996 2


3.3 Important Plots

Movies vs Users - Shows sparsity






Number of ratings given by users



Mean movie ratings given by users



4. Analysis - Model Building and Evaluation

Define RMSE: residual mean squared error

```
RMSE <- function(true_ratings, predicted_ratings){  
  sqrt(mean((true_ratings - predicted_ratings)^2))  
}
```

Model 1

Naive Mean-Baseline Model

In the first model, just based on the ratings itself, to minimize the RMSE, the best prediction of ratings for each movie will be the overall average of all ratings. The average rating is $\mu = 3.51247$, and the naive RMSE is 1.0612.

```
mu <- mean(edx$rating)  
mu
```

```
## [1] 3.51247
```

```
naive_rmse <- RMSE(validation$rating, mu)
naive_rmse
```

```
## [1] 1.0612
```

```
rmse_results <- data_frame(Model = "Just the average", RMSE = naive_rmse)
rmse_results
```


```
## # A tibble: 1 x 2
##   Model          RMSE
##   <chr>         <dbl>
## 1 Just the average 1.06120
```

Model 2

Modeling movie effects: adding `b_i` to represent average ranking for `movie_i`

Since the intrinsic features of a movie could obviously affect the ratings of a movie, we add the bias of movie/item (`b_i`) to the model, i.e., for each movie, the average of the ratings on that specific movie will have a difference from the overall average rating of all movies. We can plot the distribution of the bias and calculate the RMSE of this model.

```
movie_avgs <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = mean(rating - mu))
movie_avgs %>% qplot(b_i, geom = "histogram", bins = 10, data = ., color = I("black"))
```

```
predicted_ratings_2 <- mu + validation %>%
  left_join(movie_avgs, by='movieId') %>%
  pull(b_i)
model_2_rmse <- RMSE(validation$rating, predicted_ratings_2) # 0.943909
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Movie Effect Model",
    RMSE = model_2_rmse))
rmse_results
```

```
## # A tibble: 2 x 2
##   Model          RMSE
##   <chr>         <dbl>
## 1 Just the average 1.06120
## 2 Movie Effect Model 0.943909
```

Adding the movie bias successfully brought the RMSE to lower than 1.

Model 3

User effects: adding b_u to represent average ranking for user u

Similar to the movie effect, intrinsic features of a given user could also affect the ratings of a movie. We now further add the bias of user (b_u) to the movie effect model.


```

# mu <- mean(train_final$rating)
# just_the_sum <- train_final %>%
#   group_by(movieId) %>%
#   summarize(s = sum(rating - mu), n_i = n())
#
# rsmes[k,] <- sapply(lambdas, function(l){
#   predicted_ratings <- test_final %>%
#     left_join(just_the_sum, by='movieId') %>%
#     mutate(b_i = s/(n_i+l)) %>%
#     mutate(pred = mu + b_i) %>%
#     pull(pred)
#   return(RMSE(predicted_ratings, test_final$rating))
# })
# }
#
# rsmes
# rsmes_cv <- colMeans(rsmes)
# rsmes_cv
# qplot(lambdas, rsmes_cv)
# lambdas[which.min(rsmes_cv)] #2.2

```

We get $\lambda = 2.2$

##B) Model generation and prediction

```

lambda <- 2.2
mu <- mean(edx$rating)
movie_reg_avgs <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda), n_i = n())
predicted_ratings_4 <- validation %>%
  left_join(movie_reg_avgs, by = "movieId") %>%
  mutate(pred = mu + b_i) %>%
  pull(pred)
model_4_rmse <- RMSE(predicted_ratings_4, validation$rating) # 0.943852 not too much improved
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Regularized Movie Effect Model",
    RMSE = model_4_rmse))
rmse_results

```

```

## # A tibble: 4 x 2
##   Model                                RMSE
##   <chr>                                <dbl>
## 1 Just the average                    1.06120
## 2 Movie Effect Model                  0.943909
## 3 Movie + User Effects Model          0.865349
## 4 Regularized Movie Effect Model      0.943852

```

Model 5

Regularization of both movie and user effects (use the same lambda for both movie and user effects)

1. Perform cross validation to determine the parameter lambda

```
# define a matrix to store the results of cross validation
# lambdas <- seq(4, 8, 0.1)
# rmres_2 <- matrix(nrow=10,ncol=length(lambdas))
# # perform 10-fold cross validation to determine the optimal lambda
# for(k in 1:10) {
#   train_set <- edx[cv_splits[[k]],]
#   test_set <- edx[-cv_splits[[k]],]
#
#   # Make sure userId and movieId in test set are also in the train set
#   test_final <- test_set %>%
#     semi_join(train_set, by = "movieId") %>%
#     semi_join(train_set, by = "userId")
#
#   # Add rows removed from validation set back into edx set
#   removed <- anti_join(test_set, test_final)
#   train_final <- rbind(train_set, removed)
#
#   mu <- mean(train_final$rating)
#
#   rmres_2[k,] <- sapply(lambdas, function(l){
#     b_i <- train_final %>%
#       group_by(movieId) %>%
#       summarize(b_i = sum(rating - mu)/(n()+l))
#     b_u <- train_final %>%
#       left_join(b_i, by="movieId") %>%
#       group_by(userId) %>%
#       summarize(b_u = sum(rating - b_i - mu)/(n()+l))
#     predicted_ratings <-
#       test_final %>%
#       left_join(b_i, by = "movieId") %>%
#       left_join(b_u, by = "userId") %>%
#       mutate(pred = mu + b_i + b_u) %>%
#       pull(pred)
#     return(RMSE(predicted_ratings, test_final$rating))
#   })
# }
# rmres_2
# rmres_2_cv <- colMeans(rmres_2)
# rmres_2_cv
# qplot(lambdas,rmres_2_cv)
# lambda <- lambdas[which.min(rmres_2_cv)] #4.9
```

###2. Model generation and prediction

Regularized Movie Effect and User Effect Model

```

mu <- mean(edx$rating)
b_i_reg <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda))
b_u_reg <- edx %>%
  left_join(b_i_reg, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda))
predicted_ratings_5 <-
  validation %>%
  left_join(b_i_reg, by = "movieId") %>%
  left_join(b_u_reg, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
model_5_rmse <- RMSE(predicted_ratings_5, validation$rating) # 0.864818
rmse_results <- bind_rows(rmse_results,
  data_frame(Model="Regularized Movie + User Effect Model",
    RMSE = model_5_rmse))
rmse_results

```

```

## # A tibble: 5 x 2
##   Model                                RMSE
##   <chr>                             <dbl>
## 1 Just the average                    1.06120
## 2 Movie Effect Model                 0.943909
## 3 Movie + User Effects Model         0.865349
## 4 Regularized Movie Effect Model     0.943852
## 5 Regularized Movie + User Effect Model 0.864818

```

5. Conclusion

From the summarized RMSEs of different models, we can see that Regularization of Movie+User Model largely improved the accuracy of the prediction.

```

## # A tibble: 5 x 2
##   Model                                RMSE
##   <chr>                             <dbl>
## 1 Just the average                    1.06120
## 2 Movie Effect Model                 0.943909
## 3 Movie + User Effects Model         0.865349
## 4 Regularized Movie Effect Model     0.943852
## 5 Regularized Movie + User Effect Model 0.864818

```

The final accuracy is 0.864818